

Agentic AI & GRC Design Pattern Book

A Design Pattern Language for Policy-Driven Agentic AI Systems

By Dr. Freeman A. Jackson
Fourth Industrial Systems Corporation (4th)

© 2025 Fourth Industrial Systems Corporation

Contents

Preface	3
1 Foundational Design Patterns	4
1.1 Design Pattern 1: PolicyEnvelope	4
1.2 Design Pattern 2: Govern–Orchestrate–Retrieve–Assure (GORA)	4
2 Governance Design Patterns	7
2.1 Design Pattern 3: Policy-as-Code	7
2.2 Design Pattern 4: ConstraintEvaluator	7
2.3 Design Pattern 5: EvidenceTrail	7
3 Agentic Orchestration Design Patterns	11
3.1 Design Pattern 6: Agent Skill Registry	11
3.2 Design Pattern 7: SafeToolExecution	12
3.3 Design Pattern 8: State Machine Governance	12
4 Data, Retrieval, and RAG Design Patterns	14
4.1 Design Pattern 9: RAG Safety Envelope	14
4.2 Design Pattern 10: Knowledge Mesh	15
5 Policy and Regulatory Design Patterns	16
5.1 Design Pattern 11: Multi-Framework Crosswalk	16
5.2 Design Pattern 12: Configurable PolicyProfile	16
5.3 Design Pattern 13: Red Teaming & Safety Testing	16
6 Assurance, Testing, and Observability Patterns	19
6.1 Design Pattern 14: Compliance-Driven Test Framework	19
6.2 Design Pattern 15: Documentation Pipeline	19
A Directory Design Pattern Mapping	21
B Policy YAML Template	22
C Architecture Diagrams	23

Preface

Agentic AI systems are emerging as the next major paradigm in enterprise automation. These systems are autonomous, tool-using, multi-step, and capable of interacting with sensitive data and regulated workflows. With this power comes risk: non-determinism, hallucination, unsafe tool execution, improper data access, and difficulty demonstrating compliance to regulators.

To address this, Fourth Industrial Systems Corporation (4th) developed the **Agentic AI & GRC** framework—a multi-layer governance and orchestration system integrating:

- policy-as-code,
- agent skill registries,
- retrieval and embeddings safety envelopes,
- policy engines,
- rule evaluators,
- state-machine orchestration,
- evidence trails,
- tests and documentation pipelines.

This book formalizes the architecture into a complete **Design Pattern Language**. Each pattern encapsulates a reusable structure for safely building and governing agentic AI.

Chapter 1

Foundational Design Patterns

1.1 Design Pattern 1: PolicyEnvelope

Intent: Provide deterministic governance boundaries around non-deterministic LLM-based agents.

Problem: Agents may call unsafe tools, exceed budgets, or read sensitive data unless constrained.

Solution: Every agent action passes through a *PolicyEnvelope* containing:

- allowed/denied tools,
- risk thresholds,
- data-access tiers,
- cost governors,
- regulatory mappings,
- evidence obligations.

Repository Locations:

- `profiles/`
- `policyengine/`
- `rules/`

1.2 Design Pattern 2: Govern–Orchestrate–Retrieve–Assure (GORA)

Intent: Separate agentic AI into four explicit layers:

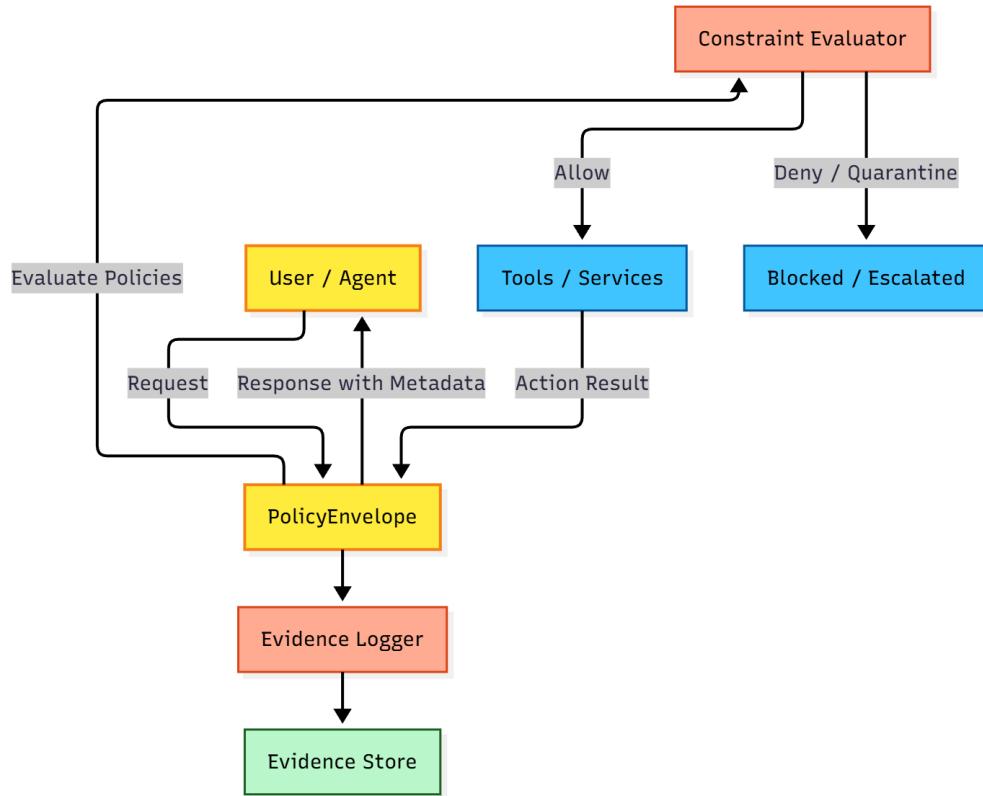


Figure 1.1: Design Pattern 1 — PolicyEnvelope diagram

1. Govern (policies, rules, decision services)
2. Orchestrate (agents, planners, workflows)
3. Retrieve (RAG pipelines, mesh search)
4. Assure (tests, monitoring, documentation)

Repository Mapping:

- Govern → `policyengine/`, `profiles/`, `rules/`
- Orchestrate → `agents/`, `functions/`
- Retrieve → `services/` (RAG, embeddings)
- Assure → `tests/`, `docs/`, `pre-commit`

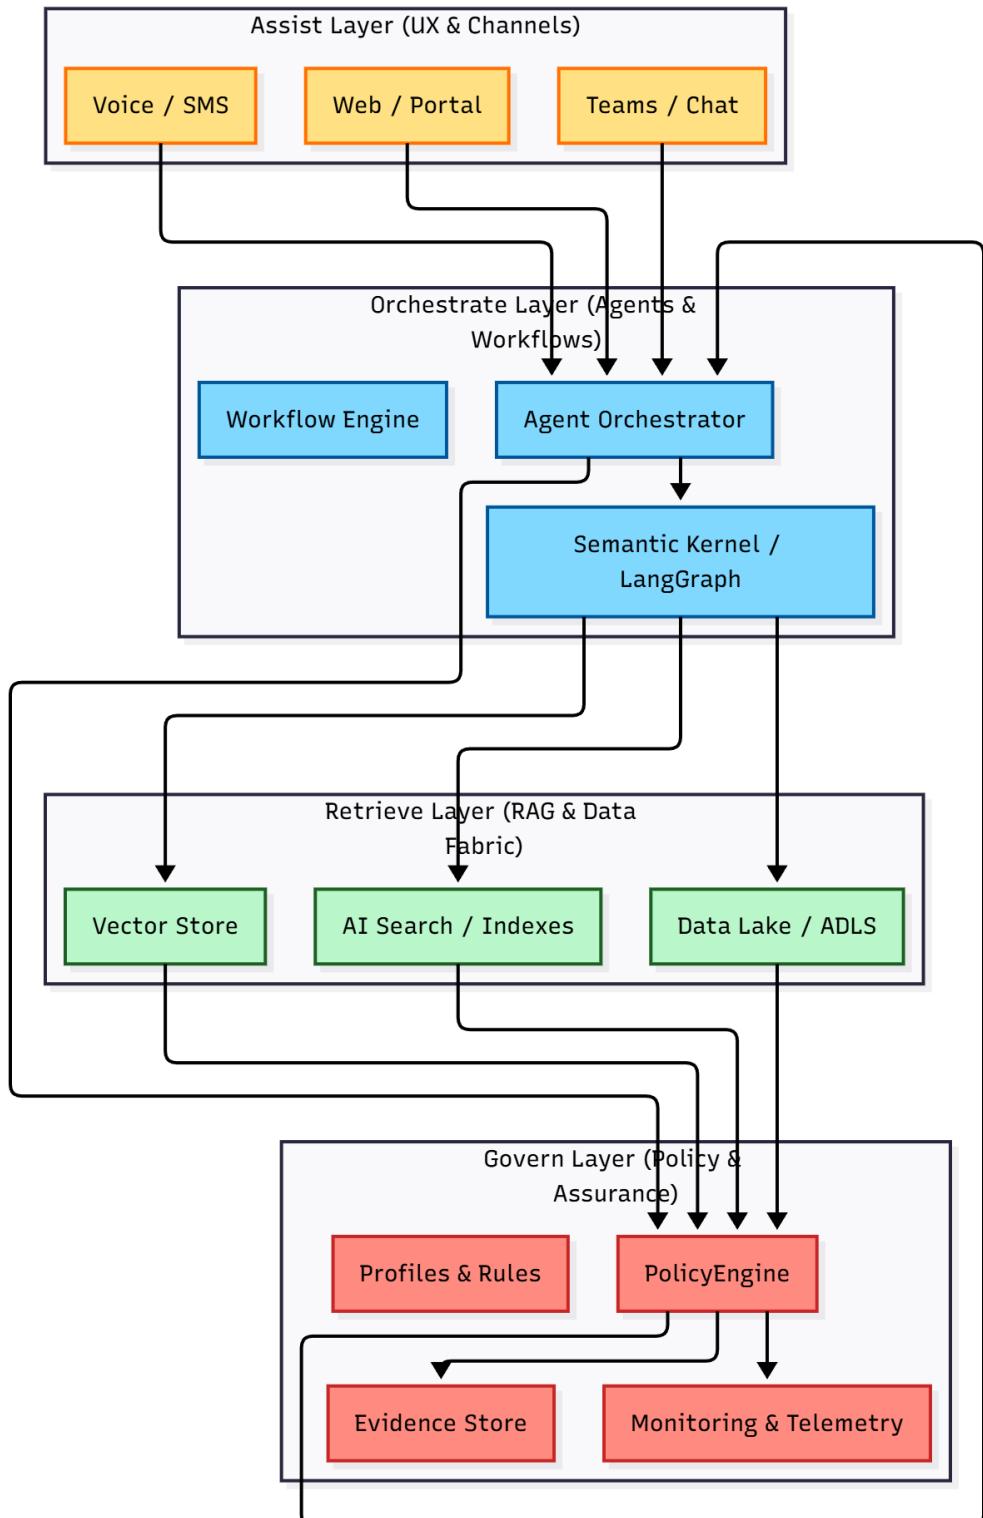


Figure 1.2: Design Pattern 2 — GORA layered architecture

Chapter 2

Governance Design Patterns

2.1 Design Pattern 3: Policy-as-Code

Intent: Express governance and safety requirements as YAML.

Listing 2.1: Example Policy Profile

```
policy:
  id: "enterprise-default"
  version: "1.0.0"
  classification: "baseline"

  allowed_tools:
    - search
    - summarizer

  denied_tools:
    - direct_db_write

  risk_thresholds:
    hallucination: 0.15
    toxicity: 0.10
    uncertainty: 0.20
```

2.2 Design Pattern 4: ConstraintEvaluator

Intent: Convert policy files into actual runtime decisions.

```
decision = evaluator.evaluate(request, profile="enterprise_default")
print(decision.outcome) # allow, deny, escalate, quarantine
```

2.3 Design Pattern 5: EvidenceTrail

Intent: Record every agent action, tool call, decision, and retrieval result.

Captured Evidence:

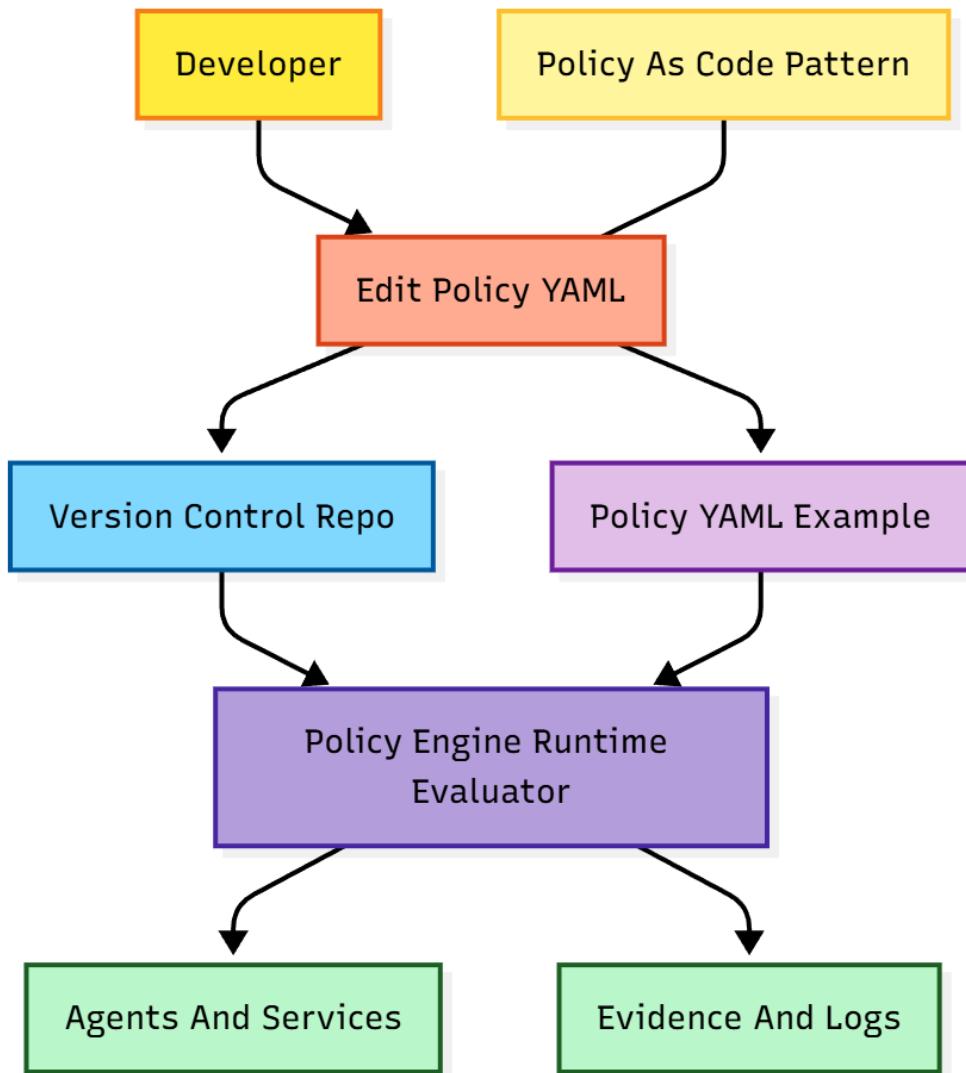


Figure 2.1: Design Pattern 3 — Policy-as-Code visual

- inputs and parameters,
- policy profile used,
- rule violations,
- RAG sources,
- model outputs,
- tool responses.

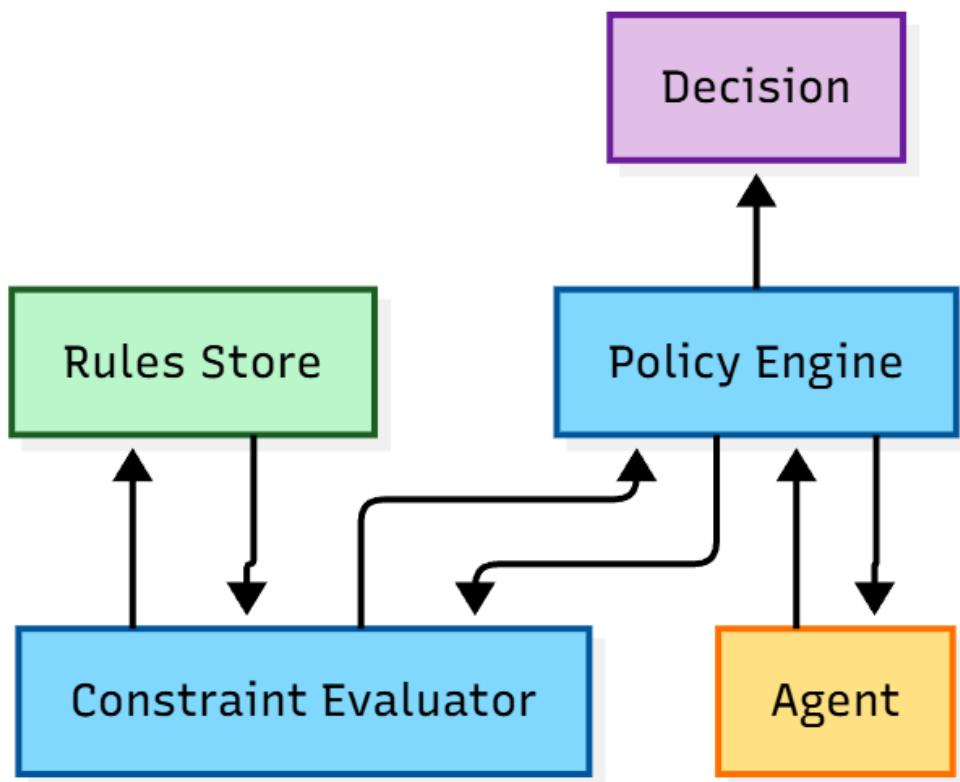


Figure 2.2: Design Pattern 4 — ConstraintEvaluator dataflow

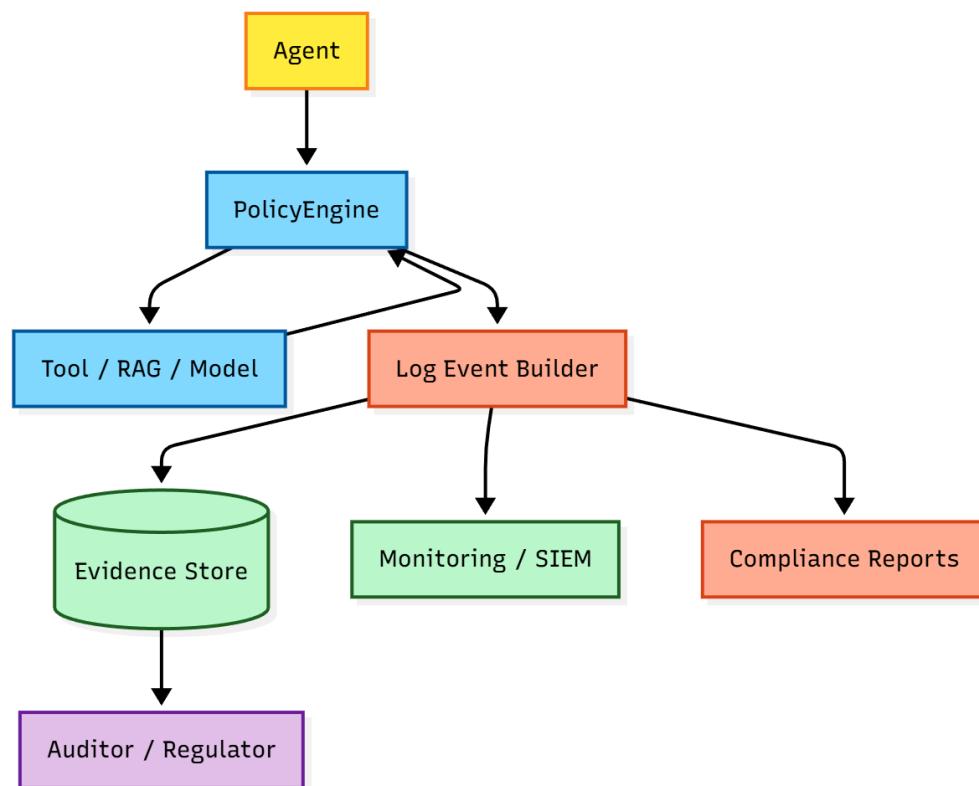


Figure 2.3: Design Pattern 5 — EvidenceTrail logging flow

Chapter 3

Agentic Orchestration Design Patterns

3.1 Design Pattern 6: Agent Skill Registry

Intent: Make agent capabilities explicit and governable.

```
SkillRegistry.register(  
    name="search_docs",  
    implementation=search_impl,  
    policy_profile="rag_restricted"  
)
```

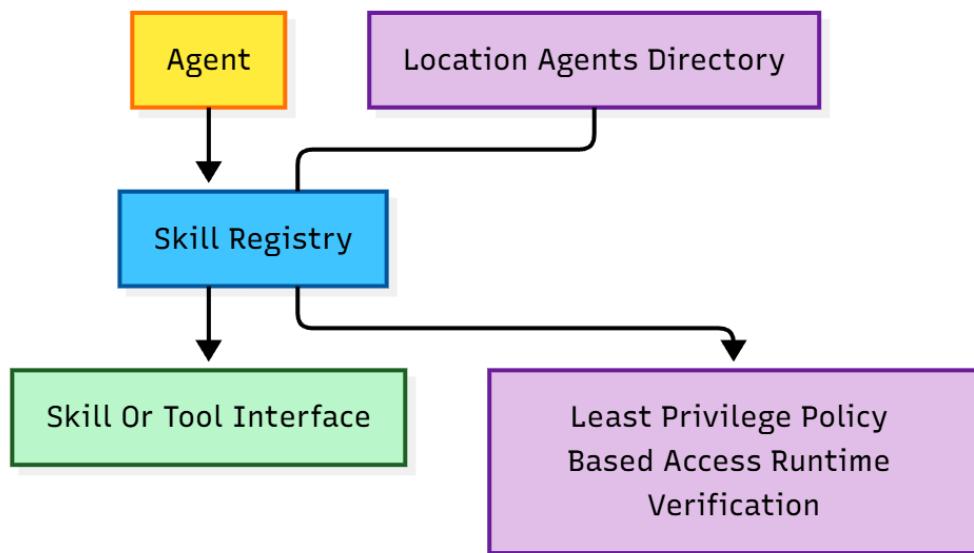


Figure 3.1: Design Pattern 6 — Agent Skill Registry

3.2 Design Pattern 7: SafeToolExecution

Intent: Ensure every tool call is checked, validated, budgeted, and logged.

```
decision = policyengine.evaluate_tool_call(agent, tool, params)
if decision.outcome != "allow":
    return make_error(decision)
```

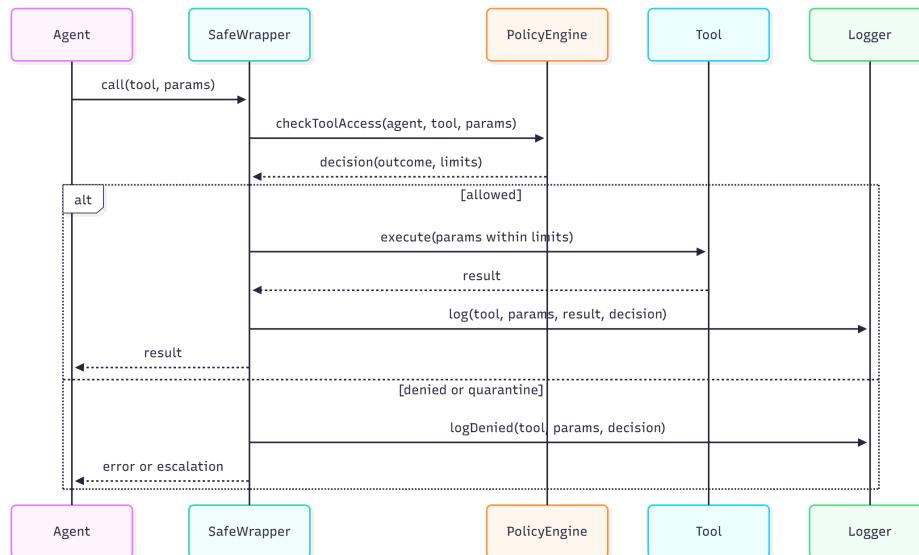


Figure 3.2: Design Pattern 7 — SafeToolExecution wrapper

3.3 Design Pattern 8: State Machine Governance

Intent: Govern multi-step workflows and transitions.

Listing 3.1: Example Workflow Definition

```

states:
  - name: "gather_requirements"
    kind: "rag_query"
    policy_profile: "rag_restricted"
  - name: "draft_response"
    kind: "llm_call"
    policy_profile: "enterprise_default"

transitions:
  - from: "gather_requirements"
    to: "draft_response"
    condition: "evidence.ok"
  
```

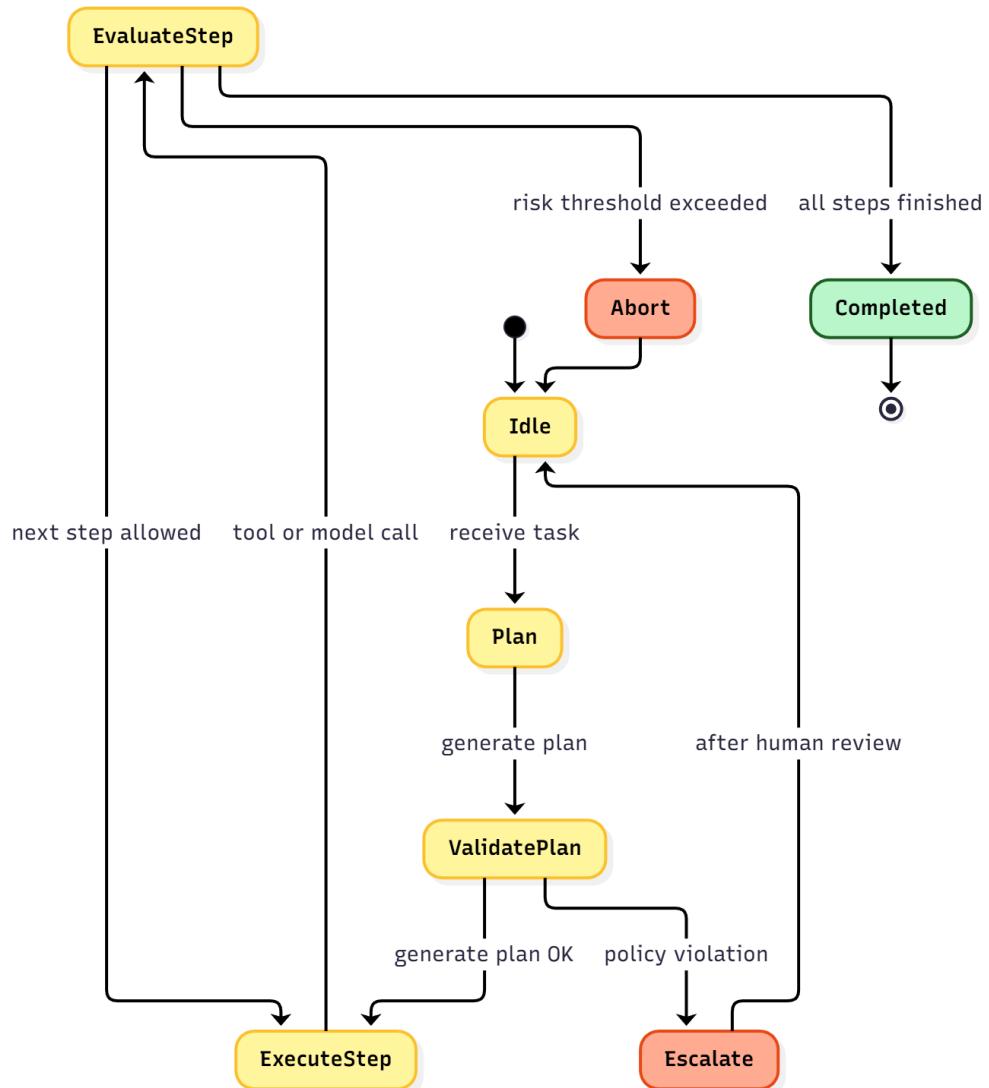


Figure 3.3: Design Pattern 8 — State Machine Governance

Chapter 4

Data, Retrieval, and RAG Design Patterns

4.1 Design Pattern 9: RAG Safety Envelope

Intent: Constrain retrieval to safe sources and safe contexts.

Key controls:

- top-k limits,
- sensitivity filters,
- domain whitelists,
- citation requirements.

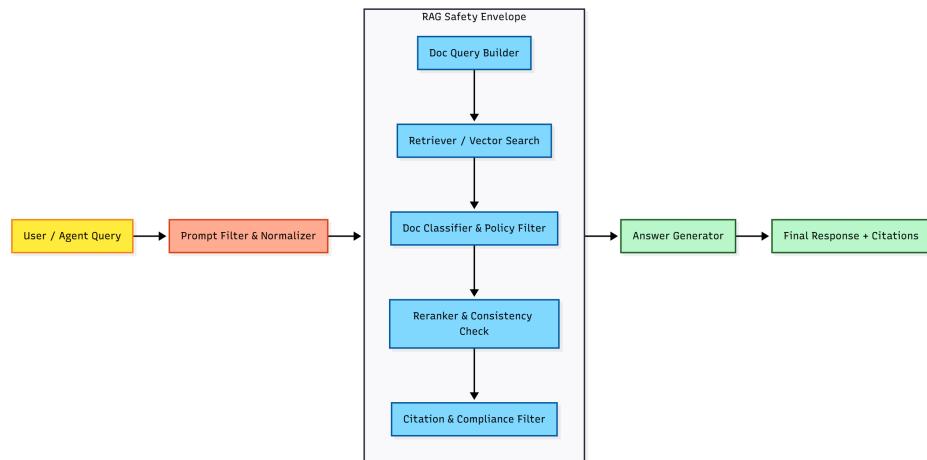


Figure 4.1: Design Pattern 9 — RAG Safety Envelope

4.2 Design Pattern 10: Knowledge Mesh

Intent: Treat retrieval sources as a governed mesh with metadata and access tiers.

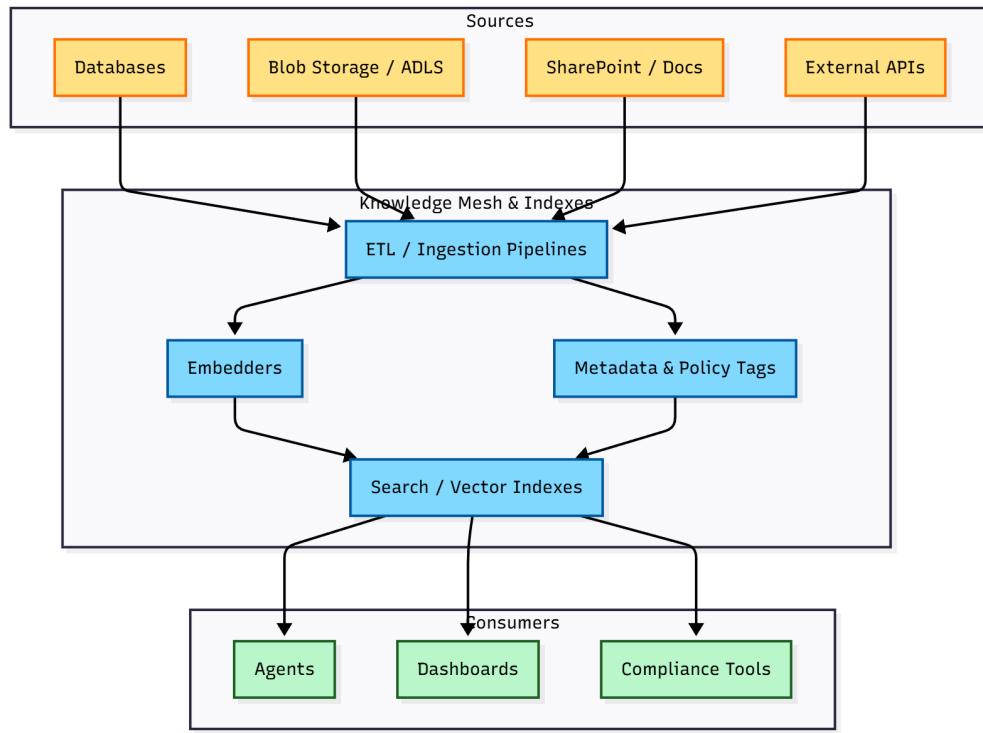


Figure 4.2: Design Pattern 10 — Knowledge Mesh topology

Chapter 5

Policy and Regulatory Design Patterns

5.1 Design Pattern 11: Multi-Framework Crosswalk

Intent: Map ISO 42001, NIST AI RMF, EU AI Act, SOC 2, etc. into profiles.

Listing 5.1: Regulatory Crosswalk Example

```
regulatory_mapping:  
  iso_42001:  
    - clause: "8.3 Traceability"  
      control: "EvidenceTrail"  
  nist_ai_rmf:  
    - function: "Govern"  
      control: "PolicyEnvelope"  
  eu_ai_act:  
    - article: "12"  
      control: "Record Keeping"
```

5.2 Design Pattern 12: Configurable PolicyProfile

Intent: Support domain-specific profiles such as:

- Healthcare (HIPAA)
- Finance (PCI)
- Defense (ITAR, CMMC)

5.3 Design Pattern 13: Red Teaming & Safety Testing

Intent: Structure red team scenarios, jailbreak prompts, and adversarial tests.

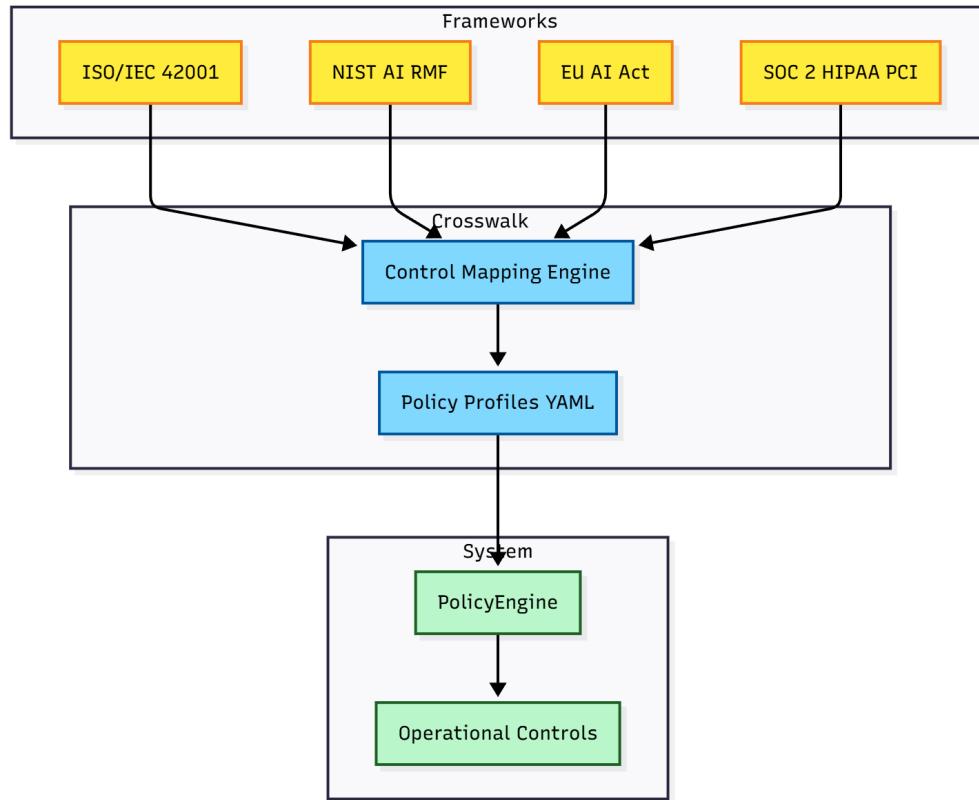


Figure 5.1: Design Pattern 11 — Multi-Framework Crosswalk mapping

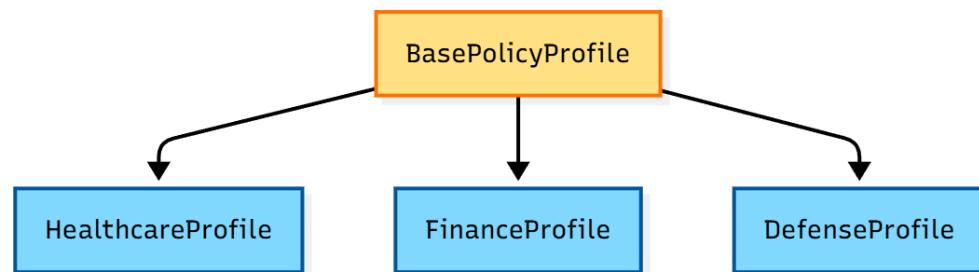


Figure 5.2: Design Pattern 12 — Configurable PolicyProfile variants

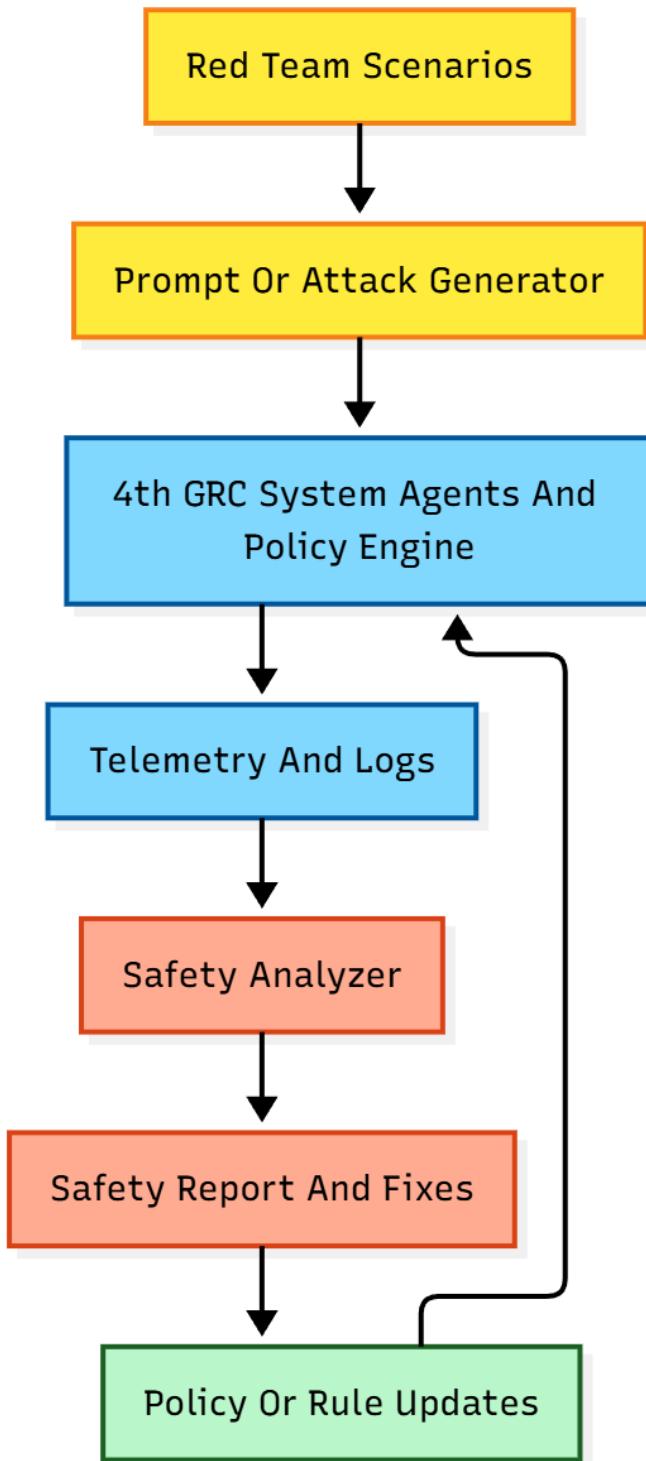


Figure 5.3: Design Pattern 13 — Red Teaming & Safety Testing workflow

Chapter 6

Assurance, Testing, and Observability Patterns

6.1 Design Pattern 14: Compliance-Driven Test Framework

Intent: Embed governance checks in CI/CD.

```
pytest tests/unit  
pytest tests/integration
```

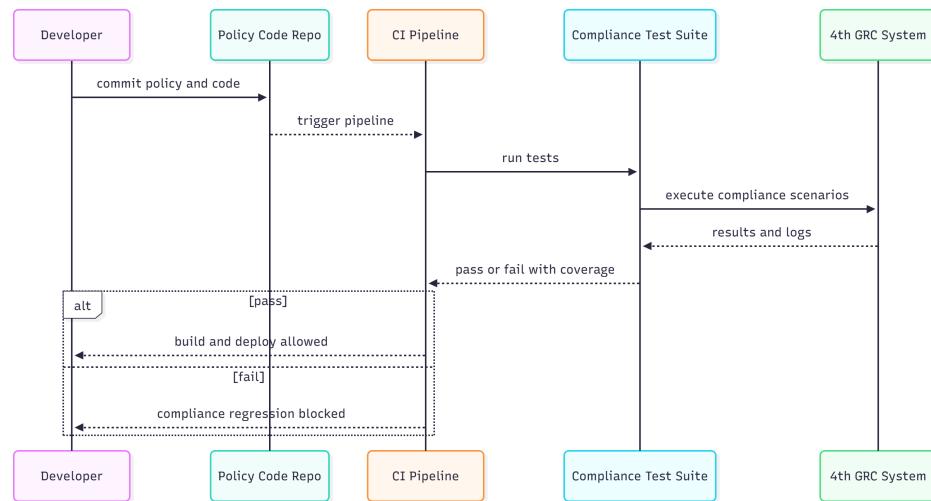


Figure 6.1: Design Pattern 14 — Compliance-Driven Test Framework pipeline

6.2 Design Pattern 15: Documentation Pipeline

Intent: Documentation (e.g., mkdocs) is treated as a governance artifact.

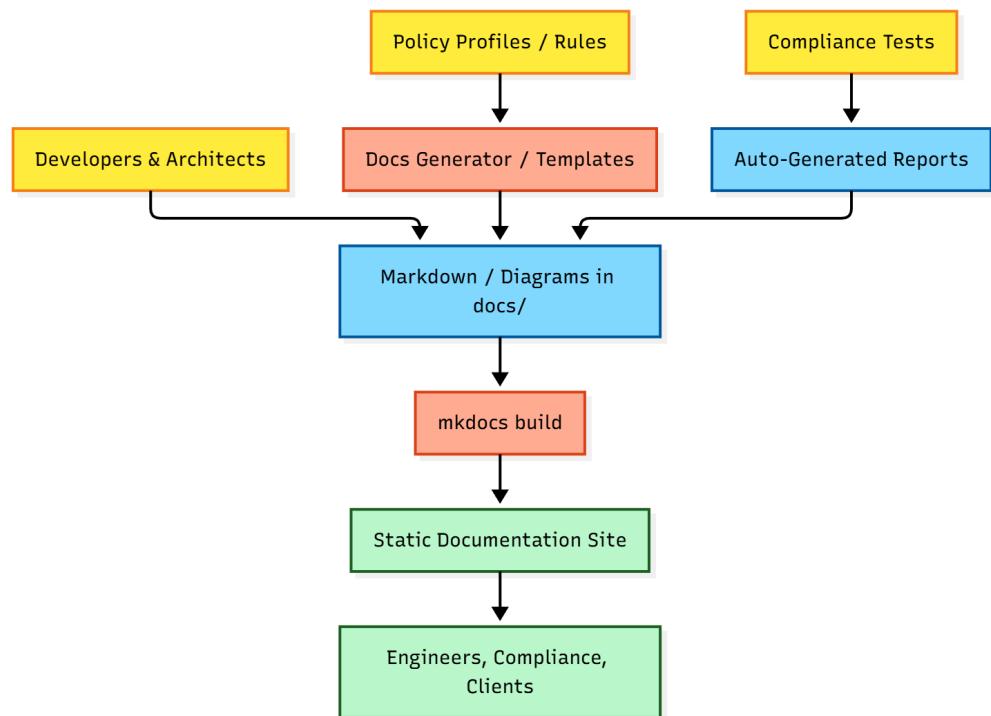


Figure 6.2: Design Pattern 15 — Documentation Pipeline flow

Appendix A

Directory Design Pattern Mapping

Directory	Design Patterns	Role
policyengine/	PolicyEnvelope, ConstraintEvaluator	Decision service for governance.
profiles/	Policy-as-Code, Configurable Profiles	YAML policies and mappings.
rules/	ConstraintEvaluator, Safety Test	Rule logic.
agents/	Skill Registry, State Machine Governance	Agent capabilities and workflows.
services/	SafeToolExecution, RAG Safety Envelope	RAG, embeddings, tool wrappers.
tests/	Assurance	Unit + integration tests.
docs/	Documentation Pipeline	System documentation.

Appendix B

Policy YAML Template

```
policy:
  id: "example-profile"
  version: "1.0.0"
  classification: "baseline"

allowed_tools:
  - search
  - summarizer

denied_tools:
  - direct_db_write
  - bulk_email

cost_governor:
  max_tokens: 4096
  max_budget: 10.00
  retry_limit: 3

data_access:
  allow:
    - PUBLIC
    - INTERNAL
  deny:
    - CONFIDENTIAL
    - RESTRICTED

risk_thresholds:
  hallucination: 0.15
  toxicity: 0.10
  uncertainty: 0.20

evidence_policy:
  capture_inputs: true
  capture_outputs: true
  store_location: "evidence/"
```

Appendix C

Architecture Diagrams

Diagrams available on request:

- C4 Model (Context, Container, Component)
- Mermaid sequence diagrams
- BPMN workflows
- SysML (Block Definition, Activity)